

Introduction to CUDA Programming

Dept. of MACS NITK



Let's learn about GPU..

A little history:

- The first GPUs were designed as graphics accelerators supported only specific fixed-function pipelines.
- In the late 1990s, the hardware became increasingly programmable
Culminating in NVIDIA's first GPU in 1999.



Graphics Processing Unit


Has thousands of cores and ALUs.

They can handle billions of repetitive low level tasks.

GPU is specialized for compute-intensive, highly parallel computation.


They are devoted to data processing rather than data caching and flow control.

Follows SPMD processing model.




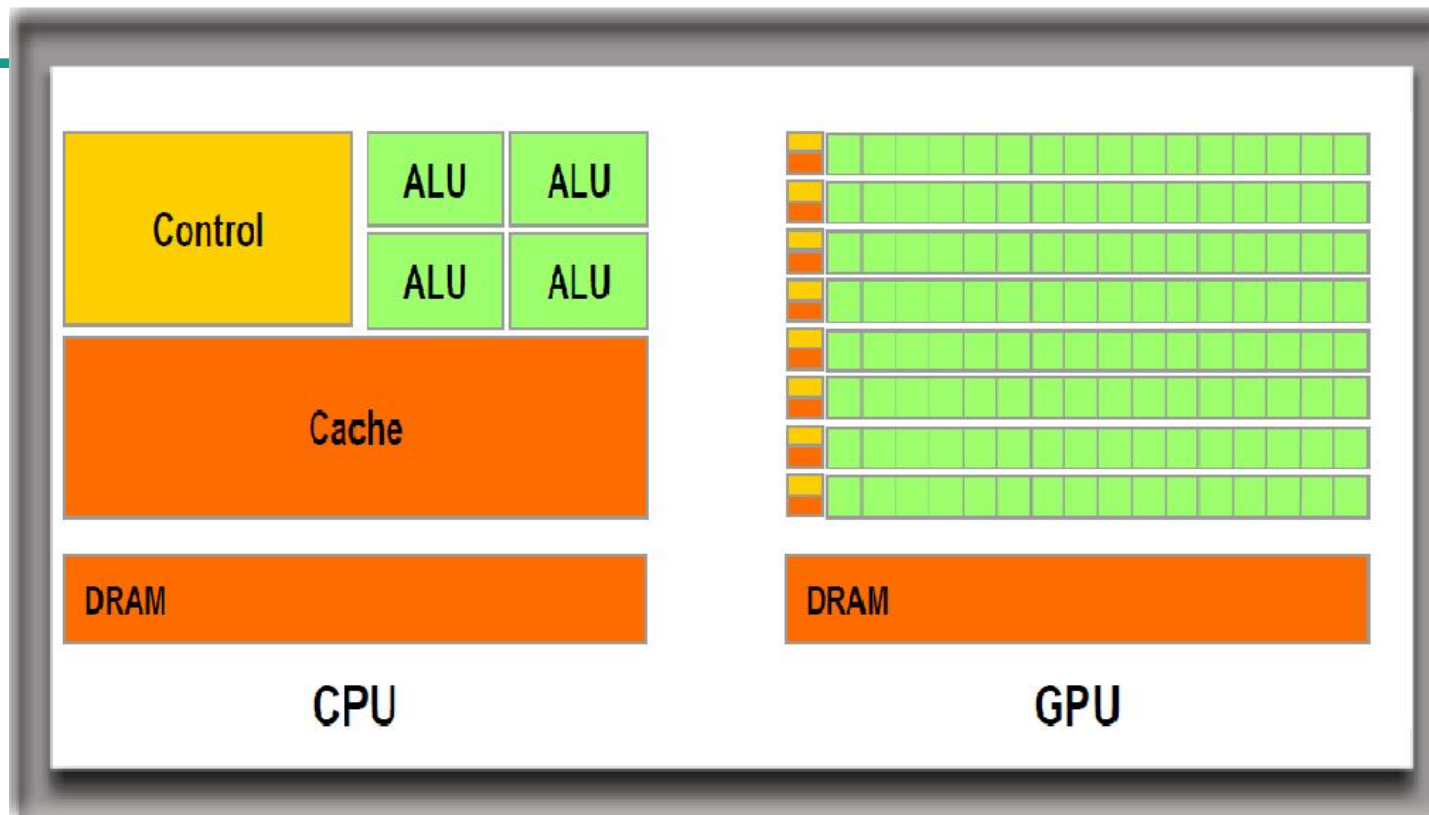


CPU (Multiple Cores)



GPU (Hundreds of Cores)







CUDA – Compute Unified Device Architecture

NVIDIA invited Ian Buck to join the company.

- Started evolving a solution to seamlessly run C on the GPU.
- Putting the software and hardware together, NVIDIA unveiled CUDA in 2006
- CUDA was launched in 2007.
- The world's first solution for general-computing on GPUs



General Structure of the GPU Program in

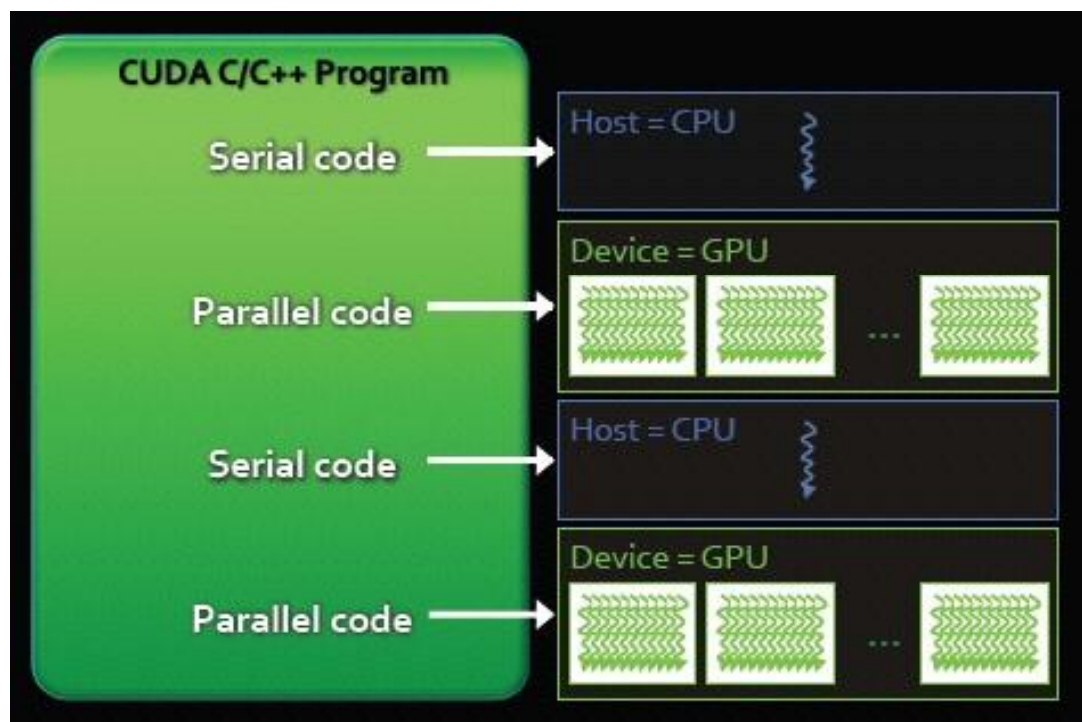
CUDA

Host Program – Executed by the CPU.

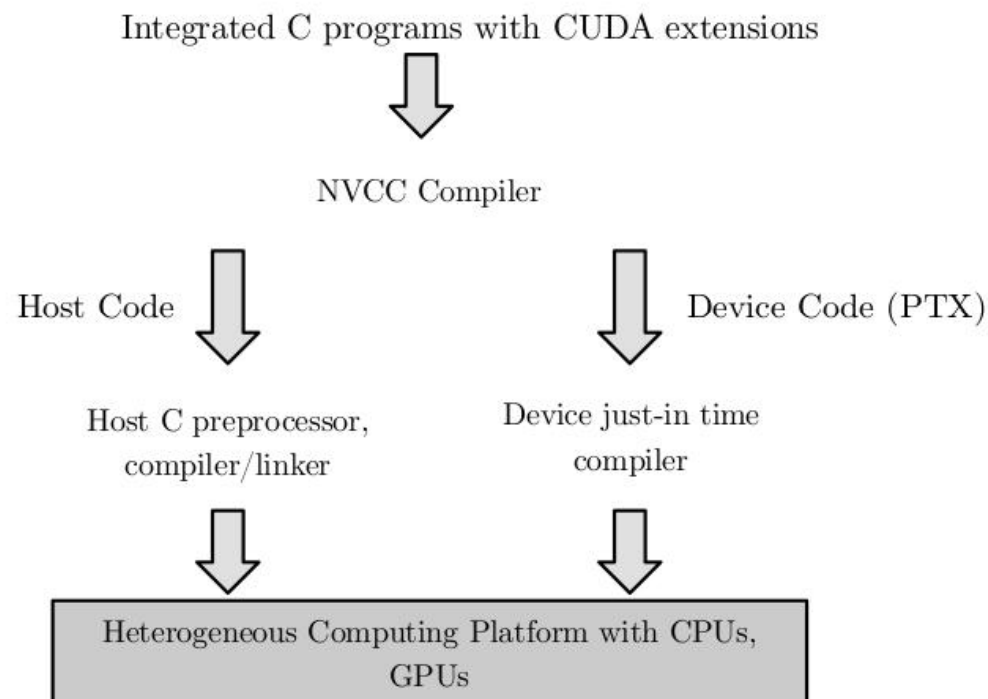
- This is a serial code.
- Sets up the parameters for GPU (kernel) execution.

Kernel Program – Executed in Parallel by the SIMD cores (Streaming Processors) in the GPU.

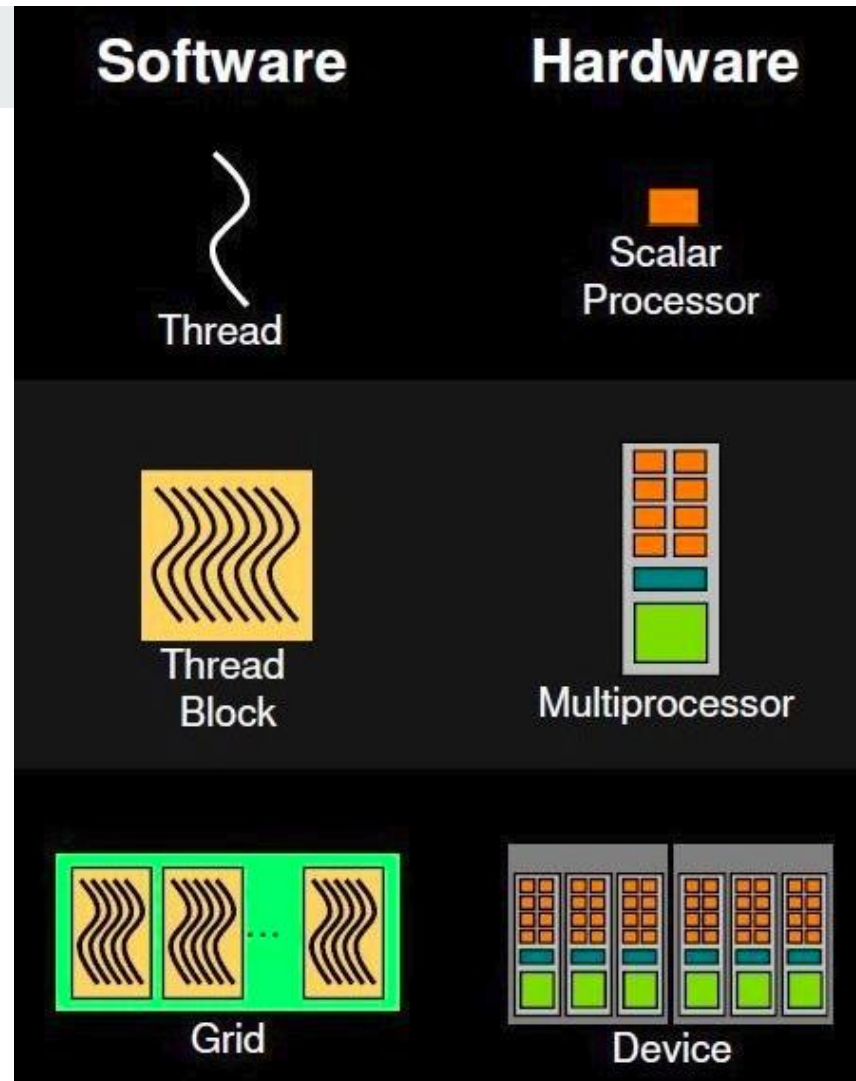
CUDA Program Execution

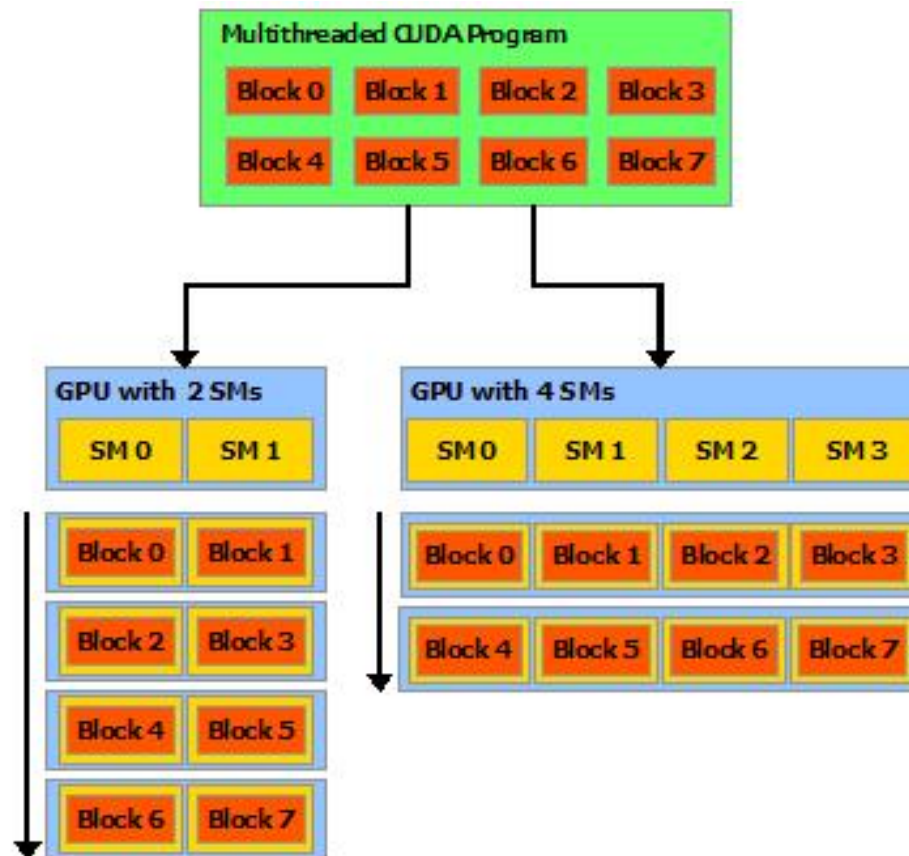


CUDA Program Execution

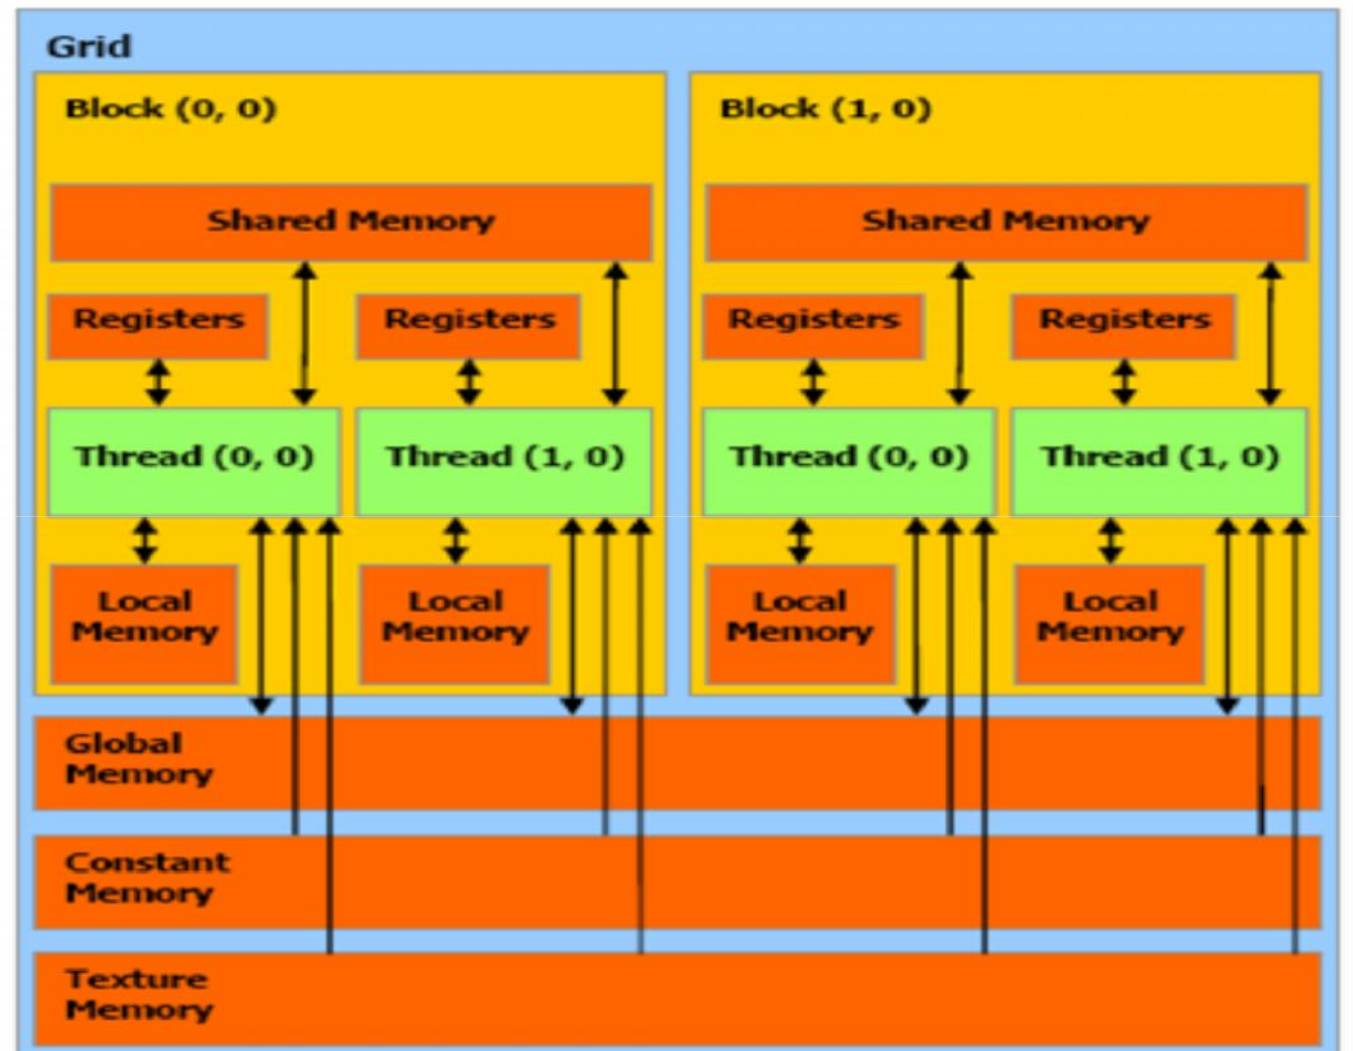


CUDA Execution Model





CUDA Memory Model:




Thread

Per-thread local
memory


Thread Block

Per-block shared
memory

Grid 0



Grid 1



Global memory



Local Memory – Accessible only by a single thread in a thread block.

Shared Memory – All threads in a single thread block have access to the shared memory.

Global Memory – All thread blocks have access to global memory. Can be **read or write**.

Constant Memory – All thread blocks have access to Constant memory. It is **read only memory**.

Texture Memory – It is a type of Global memory which is **read only cache memory**. Accessible by all the thread blocks. High speed memory.



Steps in GPU execution

Initialization of host data on CPU

Allocate memory on GPU

The data has to be transferred from CPU to GPU : Copy the data from host memory to the allocated GPU memory.

Set the kernel parameters.

Perform the computation in parallel.


Transfer back the computed result from GPU to CPU.

Free the allocated memory on GPU.




Every programming language introduction starts with Hello World!

```
int main(void)
{
    printf("Hello World!\n");
    return 0;
}
```

```
__global__ void mykernel(void) {  
  
}  
  
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!\n");  
    return 0;  
}
```




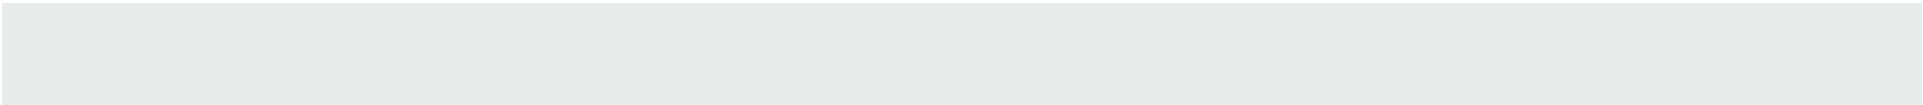
Learn about:

Keyword `__global__` indicates a function that runs on device.

It is called from host.

`mykernel()` is a device function processed by NVIDIA compiler.

Host functions are processed by host compiler.



`mykernel<<<1,1>>>()` launches the kernel

Two parameters indicate the number of blocks and number of threads per block.



Vector Addition



Vector Addition

```
// Host code
int main()
{
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...
}
```




```
// Allocate vectors in device memory
```

```
float* d_A;
```

```
cudaMalloc(&d_A, size);
```

```
float* d_B;
```

```
cudaMalloc(&d_B, size);
```


```
float* d_C;
```

```
cudaMalloc(&d_C, size);
```

```
// Copy vectors from host memory to device memory
```

```
cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
```

```
cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
```




```
// Invoke kernel
int threadsPerBlock = 256;
int blocksPerGrid =
    (N + threadsPerBlock - 1) / threadsPerBlock;
VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, N);
```

```
// Copy result from device memory to host memory
// h_C contains the result in host memory
cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);


// Free device memory
cudaFree(d_A);
cudaFree(d_B);
cudaFree(d_C);

// Free host memory
...
}
```

```
// Device code
__global__ void VecAdd(float* A, float* B, float* C, int N)
{
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```


Grid of calculus GPU= Group of blocks



gridDim.x=4
gridDim.y=3

blockIdx.x=2
blockIdx.y=1

Block (1,2)



blockDim.x=5
blockDim.y=4
blockDim.z=3

Thread (2,1,0)