



# Introduction to MPI Programming

## Dept. of MACS NITK

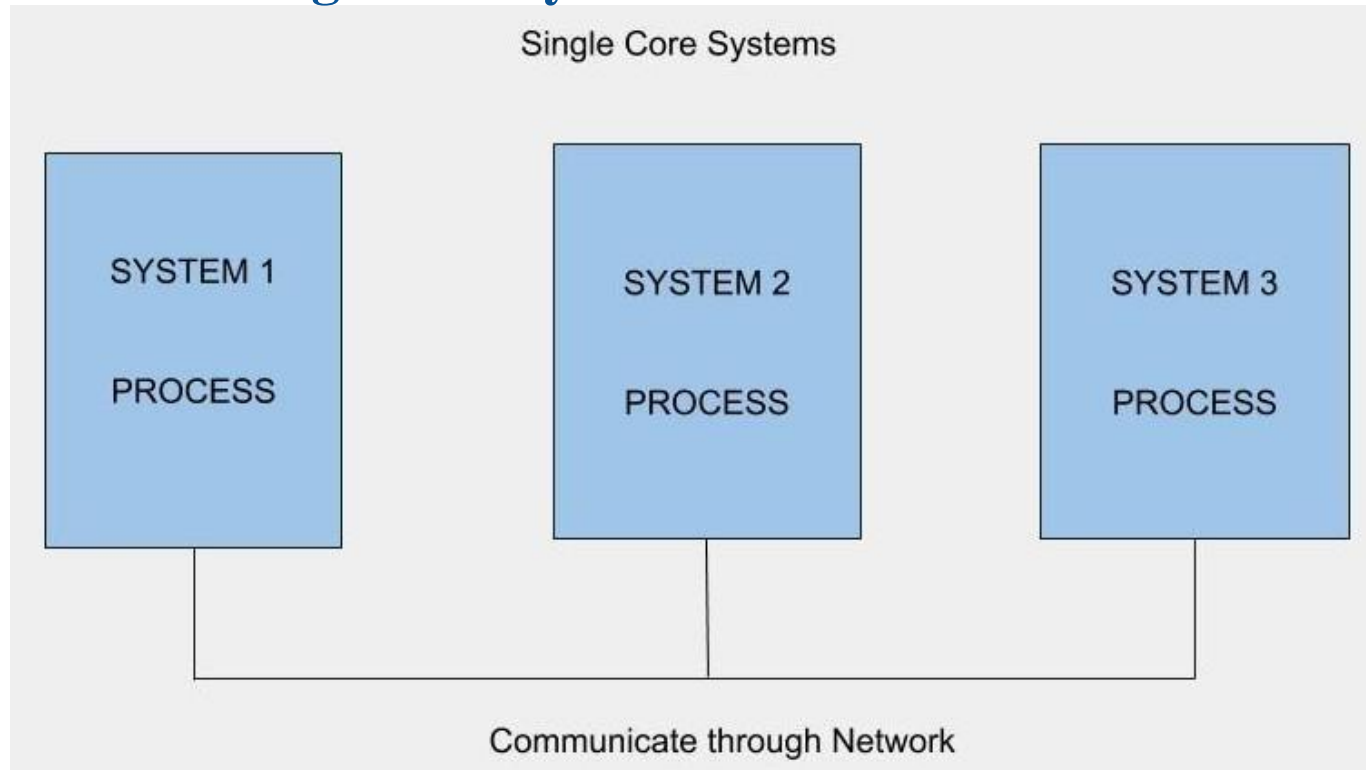
# MPI and OpenMP

- MPI – Designed for distributed memory
  - Multiple systems
    - Send/receive messages
- OpenMP – Designed for shared memory
  - Single system with multiple cores
  - Sharing memory



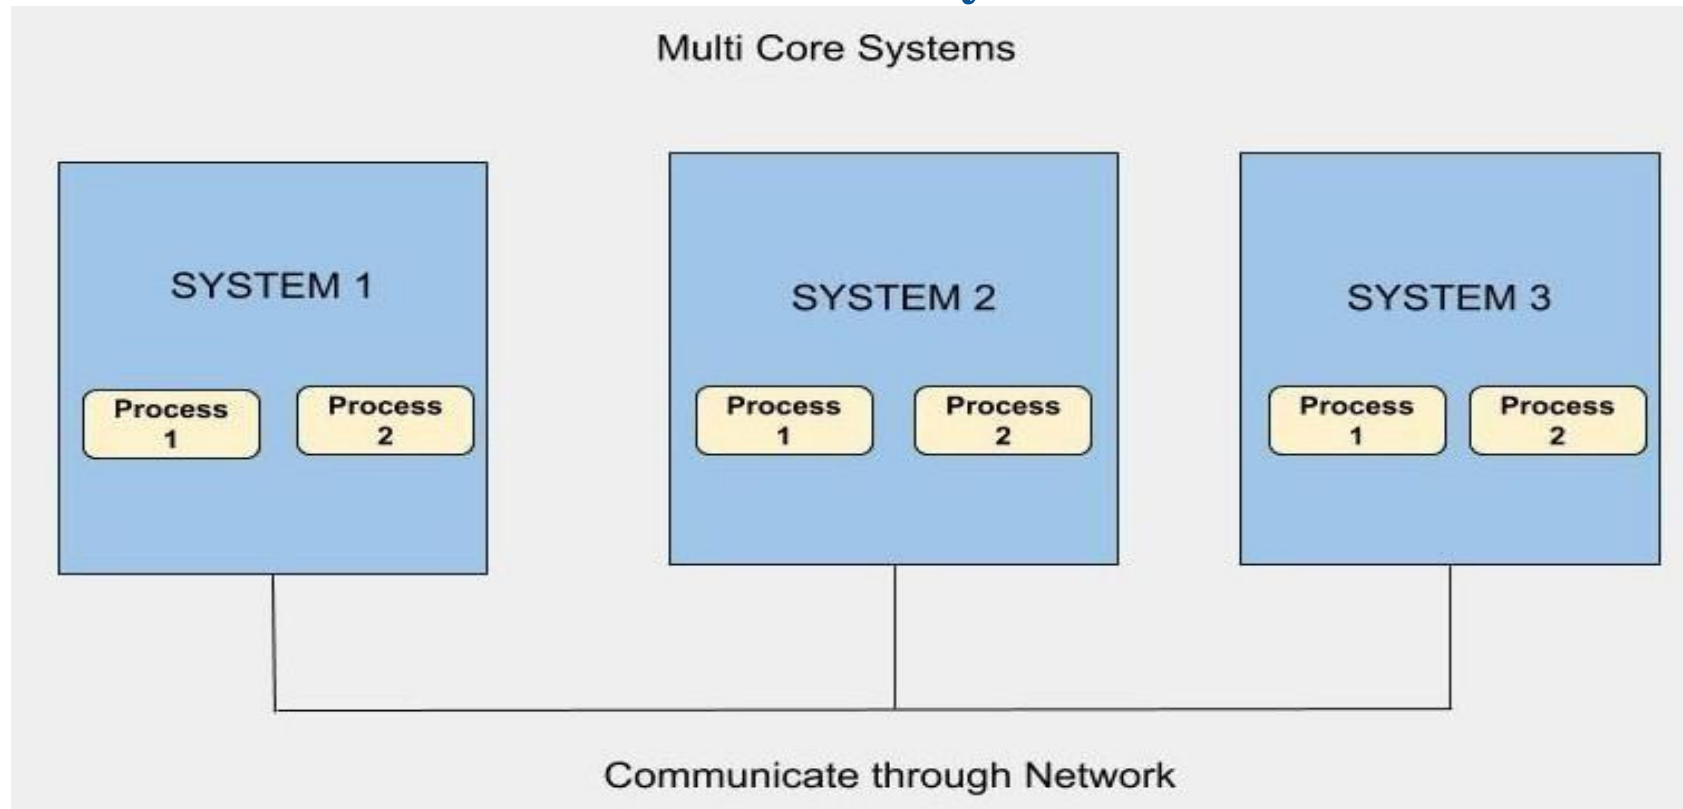
# DISTRIBUTED ARCHITECTURE

- **Systems with single core communicating through distributed memory.**
- **Heterogeneous systems**



# DISTRIBUTED ARCHITECTURE

- Systems with multiple core communicating through shared and distributed memory



# MPI

- Message Passing Interface
- Multiple implementations exist
  - Open MPI
  - MPICH
  - Many commercial (Intel, HP, C-DAC etc..)
  - Difference should only be in the compilation not development

# STRUCTURE OF MPI PROGRAM

MPI Include File

```
#include<mpi.h>
```

Initialize MPI Environment

```
MPI_Init(&argc,&argv);
```

**Computations and Message Passing**

Terminate MPI Environment

```
MPI_Finalize();
```

# Basic Environment

```
#include<stdio.h>
#include<mpi.h>
int main(int argc,char **argv)
{
    -----
    -----
    MPI_Init(&argc,&argv);
    -----
    -----
    MPI_Finalize();
    -----
    return 0;
}
```

# Communicators & Rank

- MPI uses objects called communicators
  - Defines which processes can talk
  - Communicators have a size
- **MPI\_COMM\_WORLD**
  - Predefined as ALL of the MPI Processes
  - $Size = N_{procs}$
- **Rank**
  - Integer process identifier
  - $0 \leq Rank < Size$



# Basic Environment Cont.

```
MPI_Comm_rank(comm, &rank)
```

- Returns the rank of the calling MPI process
- Within the communicator, comm
  - MPI\_COMM\_WORLD is set during Init(...)
  - Other communicators can be created if needed

```
MPI_Comm_size(comm, &size)
```

- Returns the total number of processes
- Within the communicator, comm

```
int my_rank, size;  
MPI_Init(&argc,&argv);  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &size);
```

# Hello World for MPI

```
#include <mpi.h>
#include <stdio.h>

int main (int argc, char *argv[]) {

    int rank, size;

    MPI_Init (&argc, &argv); //initialize MPI library

    MPI_Comm_size(MPI_COMM_WORLD, &size); //get number of processes
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); //get my process id

    //do something
    printf ("Hello World from rank %d\n", rank);
    if (rank == 0) printf("MPI World size = %d processes\n", size);

    MPI_Finalize(); //MPI finish

    return 0;
}
```

# To Run

Start an mpi job

With this number of processes

Run this executable

```
[user@chandra]$ mpiexec -n 4 ./hello_world_mpi
Hello World from rank 0
MPI World size = 4 processes
Hello World from rank 1
Hello World from rank 2
Hello World from rank 3
```

# Hello World Output

- 4 processes

```
Hello World from rank 3  
Hello World from rank 0  
MPI World size = 4 processes  
Hello World from rank 2  
Hello World from rank 1
```

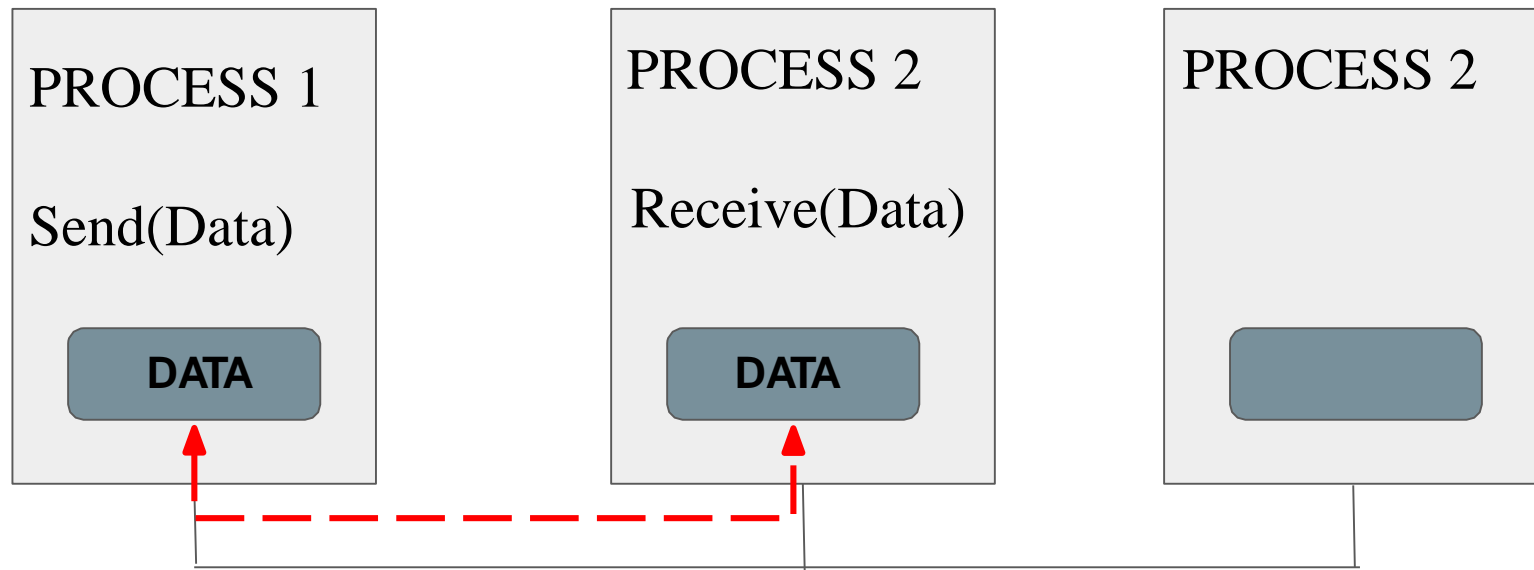
- Code ran on each process independently
- MPI Processes have *private* variables
- Processes can be on completely different machines

# Types of Message Passing:

- **Point to Point**
  - Two processes
  - Send and Receive are the basic functions
- **Collective messages**
  - Group of processes involved in communication
  - Functions like Broadcast, Scatter, Gather, Reduction

# Point to Point Communication

- Two processes involved in sending and receiving data.



- ID of sender and receiver is required.
- Specify what has to be sent and received.
- Communication needs to be synchronized.
- Communication makes use of buffers.

# Send and Receive Variants

- **Blocking Send and Receive**
- **Non Blocking Send and Receive**
- **Based on modes of Communication:**
  - **Standard**
  - **Synchronous**
  - **Buffered**
  - **Ready**

# Blocking Send and Receive

- **Basic Send and Receive routine for point to point communication.**
- **MPI Routines:**
  - MPI\_Send()
  - MPI\_Recv()



# Blocking Send and Receive

- **MPI\_Send()**

**MPI\_Send** (void \***buf**, int **count**, **MPI\_Datatype** **type**, int **dest**, int **tag**, **MPI\_Comm** **comm**)

## Parameters:

- buf** : initial address of send buffer
- count** : number of elements in send buffer (nonnegative integer)
- datatype** : datatype of each send buffer element. Ex :  
MPI\_INT, MPI\_CHAR
- dest** : rank of destination  
(integer)
- tag** : message tag (integer). For tagging send and receive.
- comm** : Communication domain of the communicating  
processes.

# Blocking Send and Receive

- MPI\_Recv():

**MPI\_Recv**(void \***buf**, int **count**, **MPI\_Datatype** **datatype**, int **source**, int **tag**, **MPI\_Comm** **comm**, MPI\_Status \***status**)

Parameters:

- |                   |  |
|-------------------|--|
| <b>buf :</b>      | initial address of receive buffer  |
| <b>count :</b>    | max number of elements in receive buffer (nonnegative integer)                                     |
| <b>datatype :</b> | datatype of each receive buffer element. Ex :<br>MPI_INT, MPI_CHAR                                 |
| <b>source :</b>   | rank of source (integer)   |
| <b>tag :</b>      | message tag (integer). For tagging send and receive.   |
| <b>comm :</b>     | Communication domain of the communicating processes.   |
| <b>status:</b>    | status object (Status). It is a structure containing information about source, tag and error code. |

# DATA TYPES

*Table 1: Basic C datatypes in MPI*

| MPI Datatype       | C datatype         |
|--------------------|--------------------|
| MPI_CHAR           | signed char        |
| MPI_SHORT          | signed short int   |
| MPI_INT            | signed int         |
| MPI_LONG           | signed long int    |
| MPI_UNSIGNED_CHAR  | unsigned char      |
| MPI_UNSIGNED_SHORT | unsigned short int |
| MPI_UNSIGNED       | unsigned int       |
| MPI_UNSIGNED_LONG  | unsigned long int  |
| MPI_FLOAT          | float              |
| MPI_DOUBLE         | double             |
| MPI_LONG_DOUBLE    | long double        |
| MPI_BYTE           |                    |
| MPI_PACKED         |                    |

# MPI Example - 1

```
for(i=0;i<50;i++) // Process 0 initializes array x
```

```
x[i]=i+1; if(myrank==0)
```

```
MPI_Send(x,10,MPI_INT,1,1,MPI_COMM_WORLD);
```

```
else if(myrank==1)
```

```
{
```

```
MPI_Recv(y,10,MPI_INT,0,1,MPI_COMM_WORLD,&status);
```

```
printf("Process %d Received Data from Process %d\n",
```

```
myrank,status.MPI_SOURCE);
```

```
for(i=0;i<10;i++)
```

```
printf("%d\t",y[i]);
```

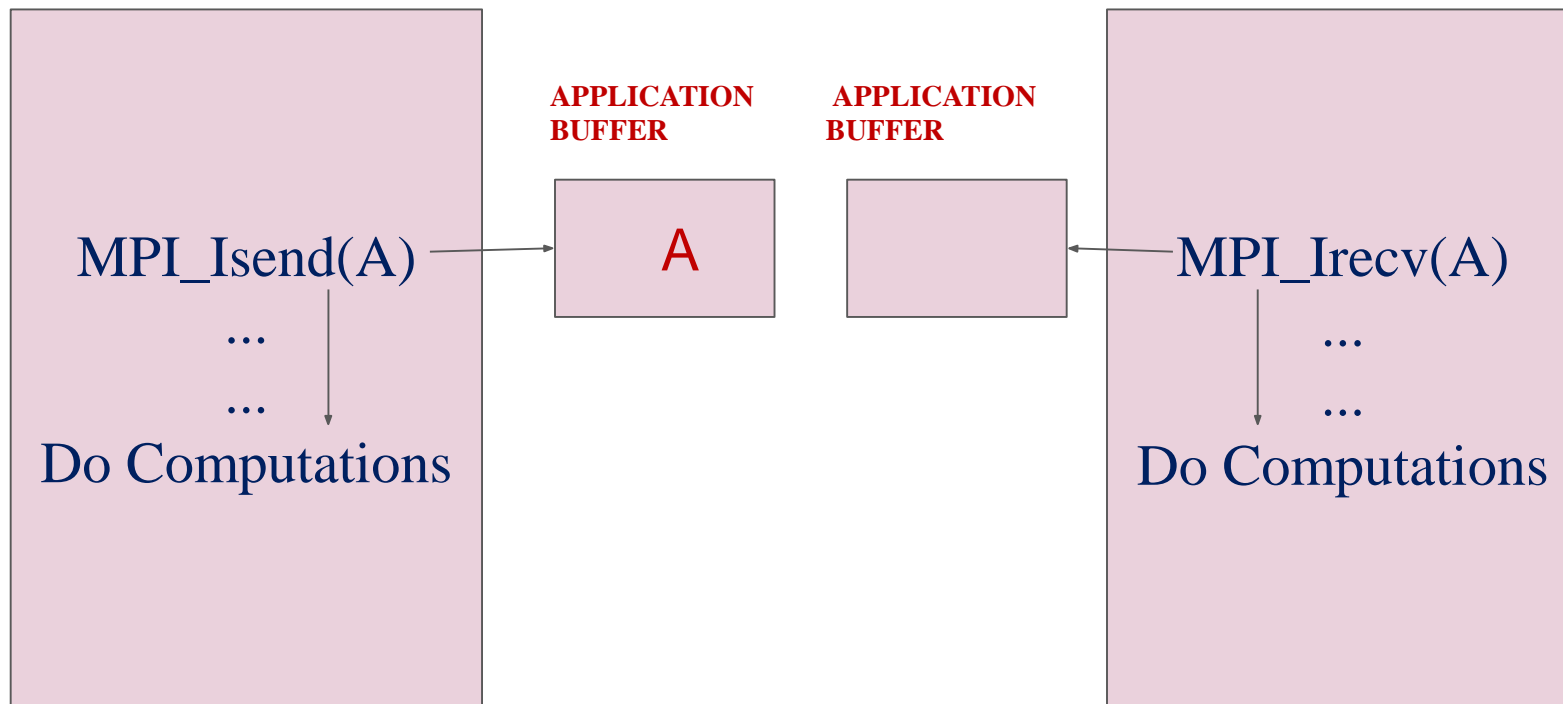
```
}
```

Process 1 Recieved data from Process 0

|   |   |   |   |   |   |   |   |   |    |
|---|---|---|---|---|---|---|---|---|----|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|---|---|---|---|---|---|---|---|---|----|

# Non Blocking Send and Receive

- Allows overlapping of computation and communication
- Advantage is Performance Gain



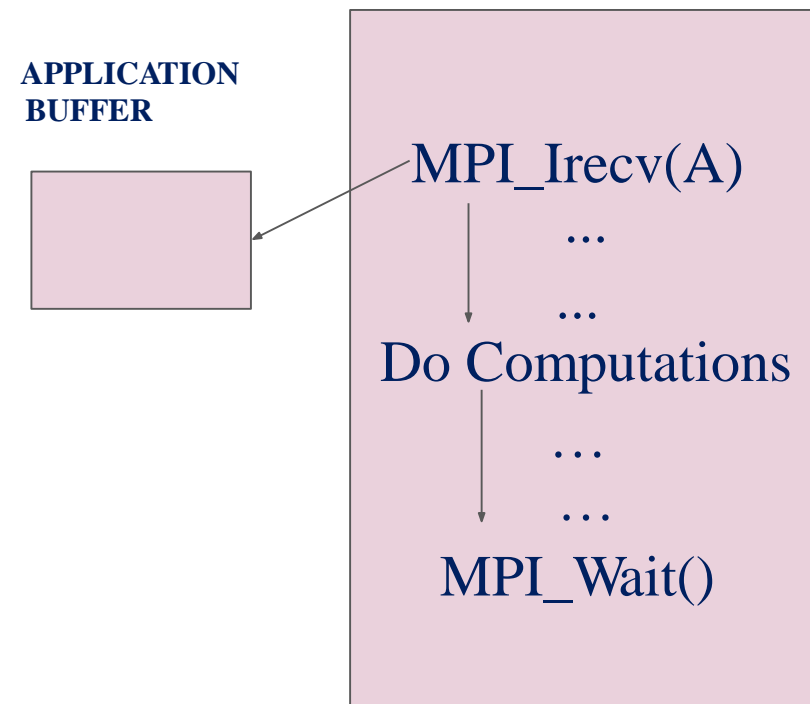
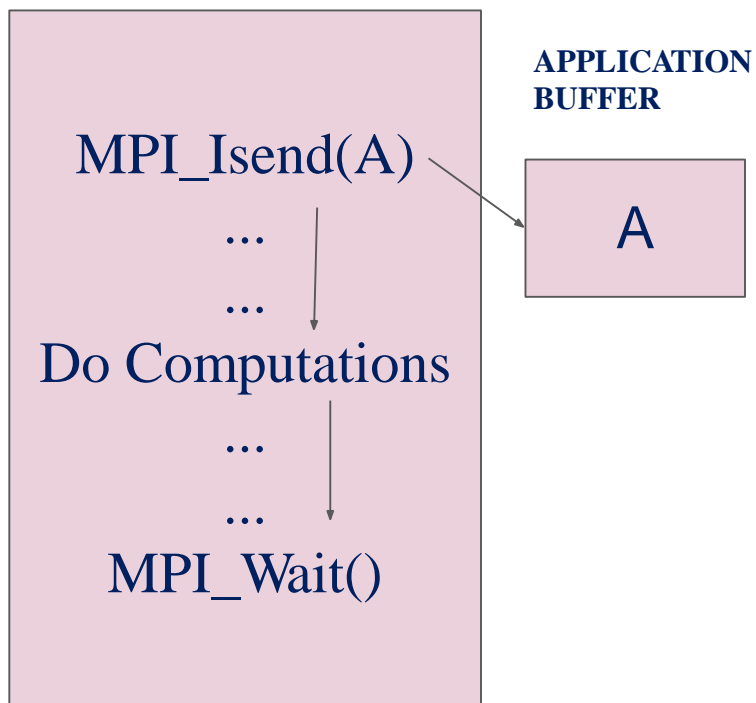
# Non Blocking Send and Receive

**MPI\_Isend** (&buf,count,datatype,dest,tag,comm,&request)

**MPI\_Irecv** (&buf,count,datatype,source,tag,comm,&request)

Parameters:

- Same as Send() and Recv() except for request
- **request** : handle. This helps to get information about MPI\_Isend and MPI\_Irecv status.
- Used in routines : **MPI\_Wait()** and **MPI\_Test()**



**Syntax :**

```
int MPI_Wait( MPI_Request *request, MPI_Status *status );
```

```
int MPI_Test( MPI_Request *request, int *flag, MPI_Status *status );
```



## MPI Example - 2

```
if(myrank==0)
{
x=10;
MPI_Isend(&x,1,MPI_INT,1,20,MPI_COMM_WORLD,&request);
printf("Send returned immediately\n");
}
else if(myrank==1)
{
MPI_Irecv(&x,1,MPI_INT,0,25,MPI_COMM_WORLD,&request);
printf("Receive returned immediately\n");
printf("Process %d of %d, Value of x is %d\n",myrank,size,x);
}
```

# What is the risk here?

```
if(myrank==0)
```

```
{ x=10;
```

```
MPI_Isend(&x,1,MPI_INT,1,20,MPI_COMM_WORLD,&request);
```

```
printf("Send returned immediately\n");
```

```
x=x+10;
```

```
}
```

**Make sure that x is available for reuse:**

```
if(myrank==0)
```

```
{ x=10;
```

```
MPI_Isend(&x,1,MPI_INT,1,20,MPI_COMM_WORLD,&request);
```

```
printf("Send returned immediately\n");
```

```
MPI_Wait(request, status)
```

```
x=x+10;
```

```
}
```

# Communication Modes

- **Standard Mode** : Calls block until message has been either transferred or copied to an internal buffer for later delivery. Ex: `MPI_Send()` and `MPI_Recv()`
- **Buffered Mode** : Send may start and return before a matching receive. `MPI_Bsend()`
- **Synchronous Mode** : Call blocks until matching receive has been posted and the message reception has started. `MPI_Ssend()`
- **Ready Mode** : Requires that a matching receive is already posted. `MPI_Rsend()`.

# MPI-Example - 3

```
if(myrank==0) {
```

```
//Blocking send will expect matching receive at the destination In Standard mode,Send will  
return after copying the data to the buffer
```

```
MPI_Send(x,10,MPI_INT,1,1,MPI_COMM_WORLD);
```

```
// This send will be initiated and matching receive is already there so the program will not lead  
to deadlock
```

```
MPI_Send(y,10,MPI_INT,1,2,MPI_COMM_WORLD);
```

```
}
```

```
else if(myrank==1)
```

```
{
```

```
//P1 will block as it has not received a matching send with tag 2
```

```
MPI_Recv(x,10,MPI_INT,0,2,MPI_COMM_WORLD,&status);
```

```
MPI_Recv(y,10,MPI_INT,0,1,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
```

```
}
```

# MPI Example - 4

```
if(myrank==0) {
```

```
    MPI_Ssend(x,10,MPI_INT,1,1,MPI_COMM_WORLD);
```

// Synchronous Blocking send will expect matching receive at the destination.  
This results in deadlock.

```
    MPI_Send(y,10,MPI_INT,1,2,MPI_COMM_WORLD); // This call will not be  
    executed
```

```
}
```

```
else if(myrank==1)
```

```
{
```

```
    MPI_Recv(x,10,MPI_INT,0,2,MPI_COMM_WORLD,&status); // P1 will block  
    as it has not received a matching send with tag 2
```

```
    MPI_Recv(y,10,MPI_INT,0,1,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
```

```
}
```


# Collective Communication

- **Broadcast**
- **Scatter**
- **Gather**
- **Reduce**
- **Scatterv**
- **Gatherv**

# Collective Communication (Bcast)

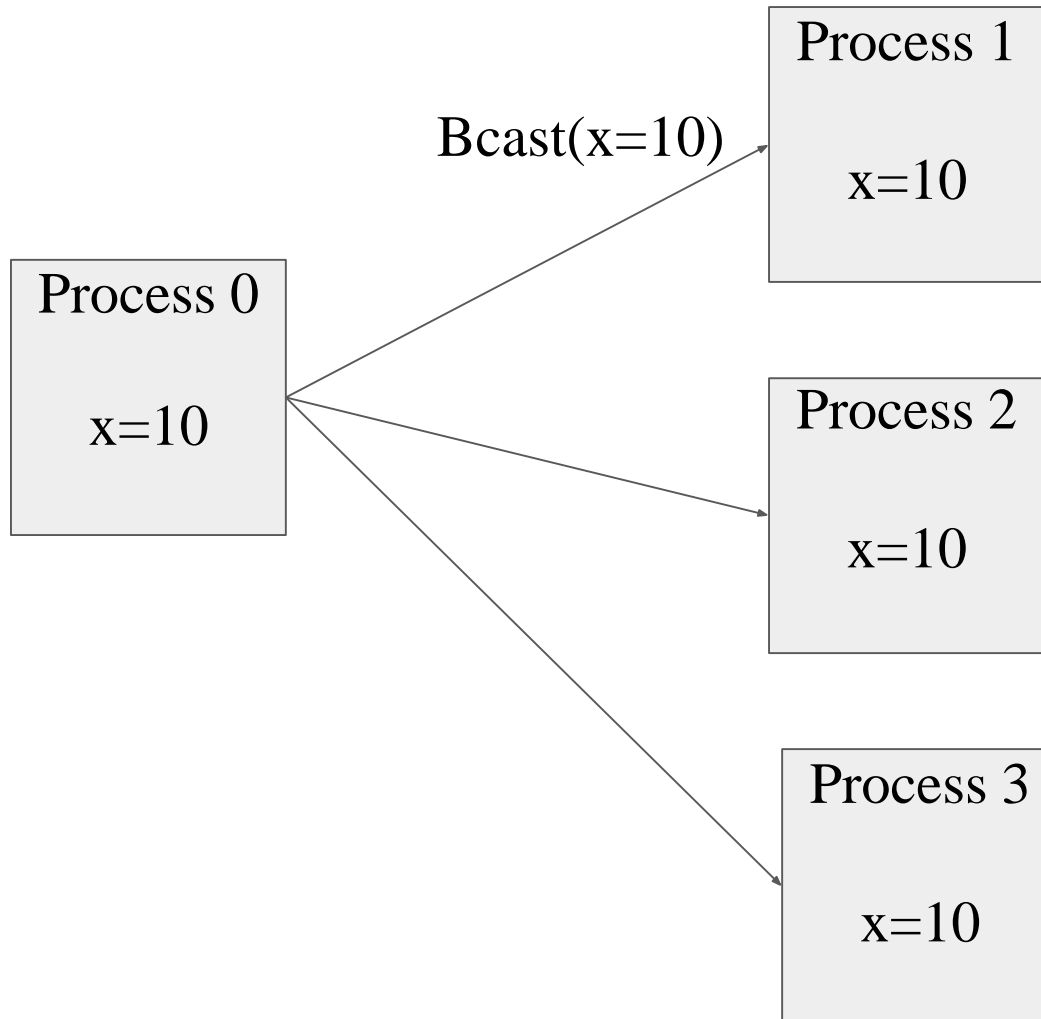
```
MPI_Bcast(&buffer, count, datatype, root, comm)
```

- Broadcasts a message from the root process to all other processes
- Useful when reading in input parameters from file





# Bcast():




## MPI Example - 5

```
if(myrank==0)
{
scanf("%d",&x);
}
MPI_Bcast(&x,1,MPI_INT,0,MPI_COMM_WORLD);
printf("Value of x in process %d : %d\n",myrank,x);
MPI_Finalize();
return 0;
}
```

# Collective Communication (Scatter)

```
MPI_Scatter(&sendbuf, sendcnt, sendtype, &recvbuf,  
           recvnt, recvtype, root, comm)
```

- Sends individual messages from the root process to all other processes




## MPI Example - 6

```
if(myrank==0)
{
printf("Enter values into array x:\n");
for(i=0;i<8;i++)
scanf("%d",&x[i]);
}
MPI_Scatter(x,2,MPI_INT,y,2,MPI_INT,0,MPI_COMM_WORLD);
for(i=0;i<2;i++)
printf("\nValue of y in process %d : %d\n",myrank,y[i]);
```

# Collective Communication (Gather)

```
MPI_Gather(&sendbuf, sendcnt, sendtype, &recvbuf,  
          recvcnt, recvtype, root, comm)
```

- Opposite of Scatter



## MPI-Example 7

**x=10, y[50]**

```
MPI_Gather(&x,1,MPI_INT,y,1,MPI_INT,0,MPI_COMM_WORLD);
```

```
// Value of x at each process is copied to array y in Process 0
```

```
if(myrank==0)
```

```
{
```

```
for(i=0;i<size;i++)
```

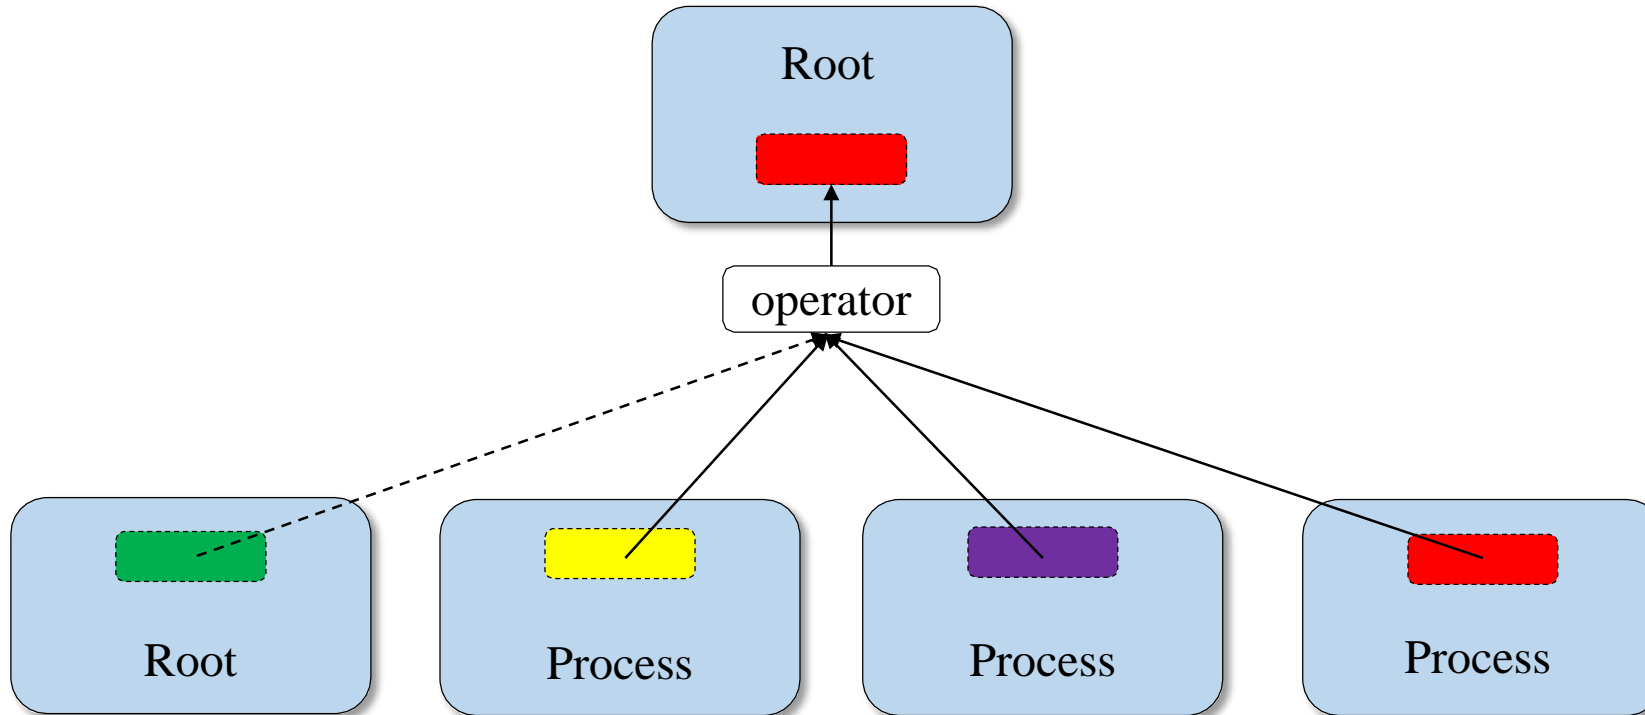
```
printf("\nValue of y[%d] in process %d : %d\n",i,myrank,y[i]);
```

```
}
```

# Collective Communication (Reduce)


```
MPI_Reduce(&sendbuf, &recvbuf, count, datatype,  
            mpi_operation, root, comm)
```

- Applies reduction operation on data from all processes
- Puts result on root process



# Collective Communication: Reduce

- Allows to perform computations on data present at multiple processes.
- Computations like : Sum, Product, Maximum, Minimum
- Stores the result in one process.





# Collective Communication: Reduce

**MPI\_Reduce**(sendbuf, recvbuf, count, datatype, operation, dest, comm)

**Parameters:**

**count:** size of receive buffer

**operation:**

| MPI name | Operation   |
|----------|-------------|
| MPI_MAX  | Maximum     |
| MPI_MIN  | Minimum     |
| MPI_SUM  | Summation   |
| MPI_PROD | Product     |
| MPI_LAND | Logical AND |
| MPI_LOR  | Logical OR  |
| MPI_LXOR | Logical XOR |

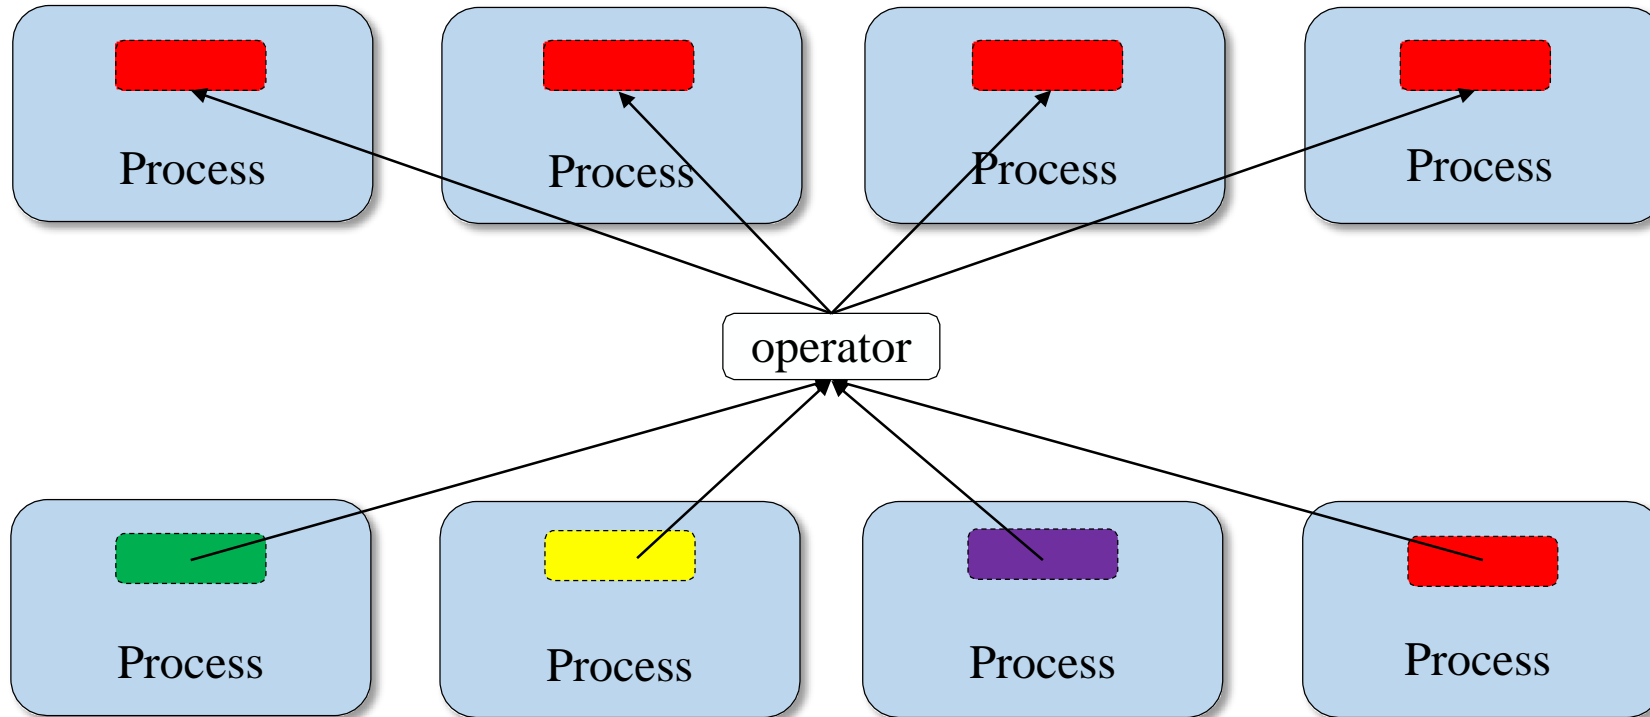
## MPI Example - 8

```
x=myrank;  
MPI_Reduce(&x,&y,1,MPI_INT,MPI_SUM,0,MPI_COMM_WORLD);  
if(myrank==0)  
{  
    printf("Value of y after reduce : %d\n",y);  
}
```

# Collective Communication (Allreduce)

```
MPI_Allreduce(&sendbuf, &recvbuf, count,  
             datatype, mpi_operation, comm)
```

- Applies reduction operation on data from all processes
- Stores results on all processes



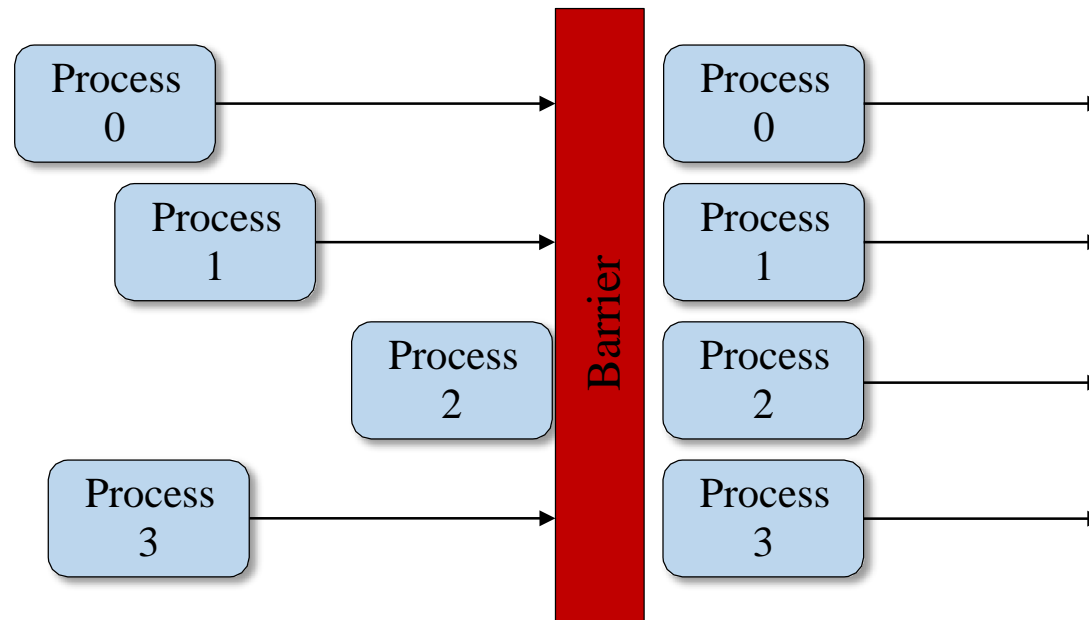
# MORE Collective Communication Routines:

- **MPI\_Gatherv()**
- **MPI\_Scatterv()**
- **MPI\_Allgather**
- **MPI\_Scan()**
- **MPI\_Barrier()**
- **MPI\_Comm\_Split()**


# Collective Communication (Barrier)

`MPI_Barrier(comm)`

- Process synchronization (blocking)
  - All processes forced to wait for each other
- Use only where necessary
  - Will reduce parallelism



# MPI\_Allgather




# MPI\_Gatherv


|   |   |   |   |
|---|---|---|---|
| 1 | 2 | 3 | 4 |
|---|---|---|---|

 counts

|   |   |   |   |
|---|---|---|---|
| 9 | 7 | 4 | 0 |
|---|---|---|---|


 displs

# MPI\_Scatterv





# MPI\_Scan



# Resources

<http://www.mpi-forum.org> (location of the MPI standard)

<http://www.llnl.gov/computing/tutorials/mpi/>

<http://www.nersc.gov/nusers/help/tutorials/mpi/intro/>

<http://www-unix.mcs.anl.gov/mpi/tutorial/gropp/talk.html>

<http://www-unix.mcs.anl.gov/mpi/tutorial/>

**MPICH** (<http://www-unix.mcs.anl.gov/mpi/mpich/>)

**Open MPI** (<http://www.open-mpi.org/>)

THANK YOU